

# IoT Intrusion Detection with Distributed Novelty Detection: Design, Implementation and Evaluation [DRAFT]

Luís Puhl, Guilherme Weigert Cassales, Hermes Senger, Helio Crestana Guardia  
Universidade Federal de São Carlos, Brasil  
Email: {luispuhl, gwcassales}@gmail.com, {hermes@, helio@dc.}ufscar.br

**Abstract**—The implementation of the Internet of Things (IoT) is sharply increasing the small devices count and variety on edge networks and, following this increase the attack opportunities for hostile agents also increases, requiring more from the network administrator role and the need for tools to detect and react to those threats. One such tool are the Network Intrusion Detection Systems (NIDS) where the network traffic is captured and analysed raising alarms when a known attack pattern or new pattern is detected. For a network security tool to operate in the context of edge and IoT it has to comply with processing time, storage space and energy requirements alongside traditional requirements for stream and network analysis like accuracy and scalability. This work addresses the construction details and evaluation of an prototype distributed IDS using MINAS Novelty Detection algorithm following up the previously defined IDSA-IoT architecture. We discuss the algorithm steps, how it can be deployed in a distributed environment, the impacts on the accuracy of MINAS and evaluate the performance and scalability using a cluster of constrained devices commonly found in IoT scenarios. We found an increase of *A* 0.0 processed network flow descriptors per core added to the cluster. Also *B* 0.0% and *C* 0.0% change in *FIScore* in the tested datasets when stream was unlimited in speed and limited to 0.0 z MB/s respectively.

**Index Terms**—novelty detection, intrusion detection, data streams, distributed system, edge computing, internet of things

## I. INTRODUCTION

The advent of Internet of Things (IoT) is growing the count and diversity of devices on edge networks, this growth increases network traffic patterns and extends opportunities for cyber attacks presenting new challenges for network administrators. To address those challenges new Network Intrusion Detection Systems (NIDS) and architectures can be explored, especially in Fog Computing and Data Stream (DS) areas.

Expected results: A system that embraces and explores the inherent distribution of fog computing in a IoT scenario opposing regular systems where data streams are collected and centralized before processing; Impact assessment of the impact of distributed, regional flow characteristics, local vs global vs distributed forgetting mechanism and other polices.

IDS characteristics and description of physical scenario.

MINAS characteristics.

Distribution and IDSA-IoT architecture.

This paper is structured as follows: Section II presents previous works that addresses related problems and how they influenced our solution. Section IV address our proposal, the work done, issues found during implementation and discusses parameters and configurations options and how we arrived at our choices. Section V shows experiments layouts and results, we compare serial and distributed implementation's metrics for validation, we also evaluate communication delay effects on classification metrics and conclude with the speedup per core and overall maximum stream speed. Section VI summarizes the research results and presents our final conclusions and future works.

## II. RELATED WORK

Recent works explored those areas, to name a few: BigFlow [1] employing Apache Kafka and Apache Flink for distributed stream processing evaluating with package stream dataset, CATRACA [2], [3] uses Apache Kafka and Apache Spark for stream processing and

## III. PROPOSAL

Fog computing infrastructure aims to offload computing resources from cloud providers by placing edge devices closer to end-users and/or data sources.

Objective: Distributed novelty detection in streams using limited hardware.

Previous attempts to attain the objective of distributed fast

## IV. IMPLEMENTATION

The original MINAS algorithm has a companion implementation (*Ref*) written in Java using MOA library base algorithms such as K-means and CluStream. *Ref* employs Java's double, a 64bits number whose precision is not absolutely necessary and, as it is often necessary to shuffle between nodes via network and a small economy could be made with only a float number with 32bits. Another difference between *Ref* and *MFOG* is cluster radius calculation from the distances of elements forming the cluster and the cluster's center, where the former uses the maximum distance, the latter uses the standard deviation of all distances as described in [4].

The stream format for input and output also of note. Input information needed is the value of the item, this value is a

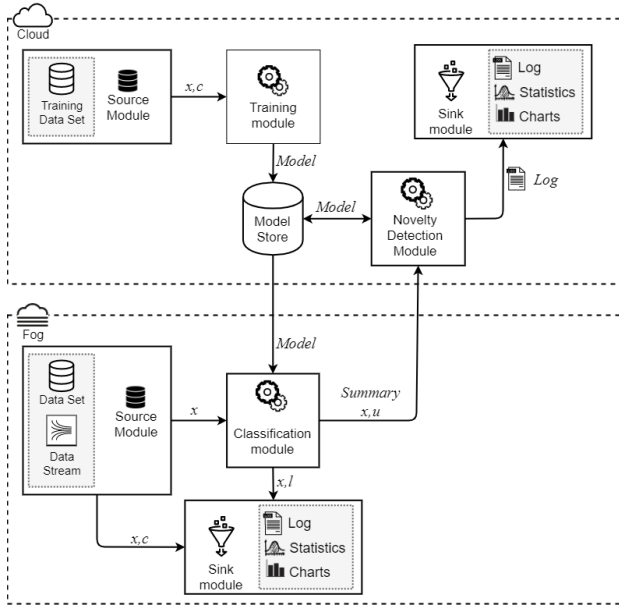


Fig. 1. MFOG architecture overview.

number sequence of length  $d$  (referenced as dimension). In addition to the value for evaluation and training purposes the class identifier as single character, optimality an unique item identifier ( $uid$ ) can be provided. For output information and format the decision isn't so clear as we can't predict future system integrations needs like only novelty alarms or every item's original value with assigned label so, we have a compromise and put only enough information for the Evaluation Module (where the full information from the testing file or stream can be accessed) meaning the format can be defined as a tuple containing  $uid$  and assigned label.

Adicionar fig arch fisica.

Another implementation decision related to the output stream is whether or not to reprocess, and add to the output stream, examples in the unknown buffer after the novelty detection procedure, meaning one item can be classified once as unknown and again with a label. Our tests using this technique had increased true positives when compared to not using it. However this changes the stream operator behavior from a *Map* to a *FlatMap* having duplicate entries on the output stream as previously mentioned. Regardless of choice the classification of the unknown buffer after a model update, using the full model or just the added set of clusters, is done to remove the examples "consumed" in the creation of a new cluster in the internals of the clustering algorithm.

For MFOG the Message Passing Interface (MPI) library was used. In MPI programming, multiple processes of the same program are created by the runtime and each process instance receives a rank parameter, for MFOG this parameters indicate if the process is root, rank 0, or leaf otherwise. Beyond this division, each process also operates two threads, for the root there is a sampler and detector threads, for the leafs each has a model receiver thread and multiple classifier threads. The

overall sequence of creation is shown in Figure 2.

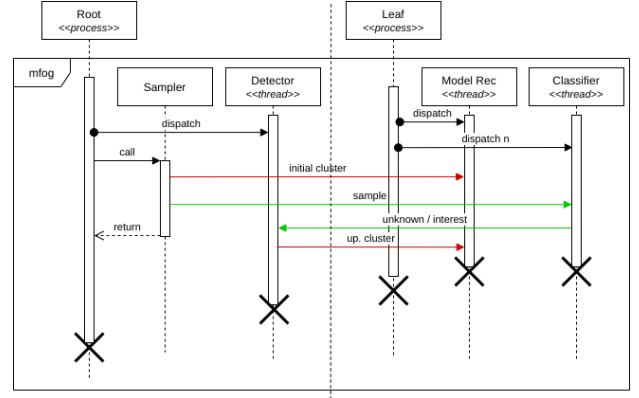


Fig. 2. MFOG life line overview.

### Polices

- Detecção de novidades e manutenção de modelo em ambiente distribuído;
- Mecanismo de ND local (síncrono) vs nuvem quanto à atraso de definição de modelo (nesse ponto é onde a hipótese prevê maior diferença, grande ponto de interesse);
- Mecanismo de esquecimento local vs global (modelo único ou por nó);
- Atraso na reclassificação dos desconhecidos;

The Evaluation Module was also build following reference techniques like multi-class confusion matrix with label-class association [4] to extract classification quality metrics.

## V. EXPERIMENTS AND RESULTS

For the experimental setup we dedicated 3 Raspberry Pi single board computers connected via Gigabit Ethernet Switch forming a simple cluster. This cluster stored all source code, binaries (compiled and linked in place) and datasets, being accessed via our laboratory network over Secure Shell (SSH). All experiments were executed in this cluster for isolation of otherwise unforeseen variations.

The dataset used is the December 2015 segment of Kyoto 2006+ Dataset<sup>1</sup> (Traffic Data from Kyoto University's Honeypots) [5]. This segment was filtered (from 7,865,245 instances) to only examples associated to known attack types identified by existing IDS, and attack types with more than 10,000 instances for significance as done by [6]. The remaining instances then were transformed by normalization, transforming each feature value space (e.g. IP Address, Duration, Service) to the Real interval  $[0,1]$ . The result is stored in two sets, training set and test set, using the holdout technique filtering in only normal class resulting in 72,000 instances for training set and 653,457 for test set, containing

O que quer testar com os experimentos. Tese: Mostrar que detecção por novidade e classificação continua viável em fog. Seria inviável por conta do atraso de distribuição de modelo

<sup>1</sup>Available at [http://www.takakura.com/Kyoto\\_data/](http://www.takakura.com/Kyoto_data/)

[illegible][illegible]

Como realizou (cenário, rpi, setup, coleta de métricas).  
Quais resultados obteve.

## Como interpretar os resultados.

The graph displays the performance of three models on the Minas Ref. mk1.1 dataset. The blue line represents the 'mis' model, the green line represents the 'hit' model, and the orange line represents the 'unk' model. The 'mis' model shows the highest precision, followed by the 'hit' model, and the 'unk' model shows the lowest precision. The x-axis is labeled with recall values from 0 to 600k, and the y-axis is labeled with precision values from 0.0 to 1.0. Vertical dashed lines are drawn at 5k, 10k, and 12k recall.

Model	Precision
mis	0.676
hit	0.306
unk	0.018

## VI. CONCLUSION

## ACKNOWLEDGMENT

Serial Online-ND

Iteration (k)	mis	hit	unk
0	0.65	0.35	0.031
56	0.65	0.35	0.031
500	0.65	0.35	0.031
600	0.659	0.310	0.031

## REFERENCES

- TABLE III

[illegible]

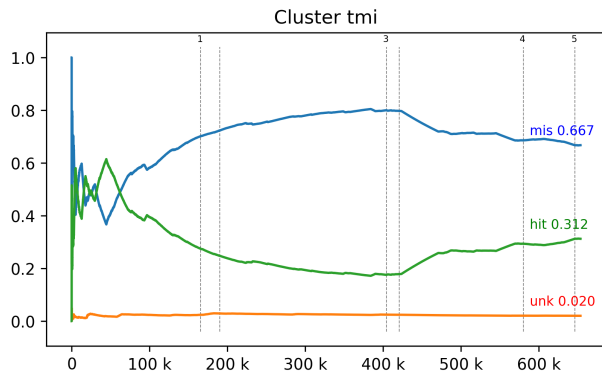


Fig. 5. Parallel implementation: Stream hits and novelties visualization

Using Stream Processing as a Virtual Function for Big Data,” Ph.D. dissertation, 2019.

- [4] E. R. d. Faria, A. C. Ponce de Leon Ferreira Carvalho, and J. Gama, “Minas: multiclass learning algorithm for novelty detection in data streams,” *Data Mining and Knowledge Discovery*, vol. 30, no. 3, pp. 640–680, May 2015. [Online]. Available: <https://doi.org/10.1007/s10618-015-0433-y>
- [5] J. Song, H. Takakura, Y. Okabe, M. Eto, D. Inoue, and K. Nakao, “Statistical analysis of honeypot data and building of Kyoto 2006+ dataset for NIDS evaluation,” *Proceedings of the 1st Workshop on Building Analysis Datasets and Gathering Experience Returns for Security, BADGERS 2011*, pp. 29–36, 2011.
- [6] G. W. Cassales, H. Senger, E. R. DE FARIA, and A. Bifet, “IDSA-IoT: An Intrusion Detection System Architecture for IoT Networks,” in *2019 IEEE Symposium on Computers and Communications (ISCC)*, June 2019, pp. 1–7. [Online]. Available: <https://ieeexplore.ieee.org/document/8969609/>